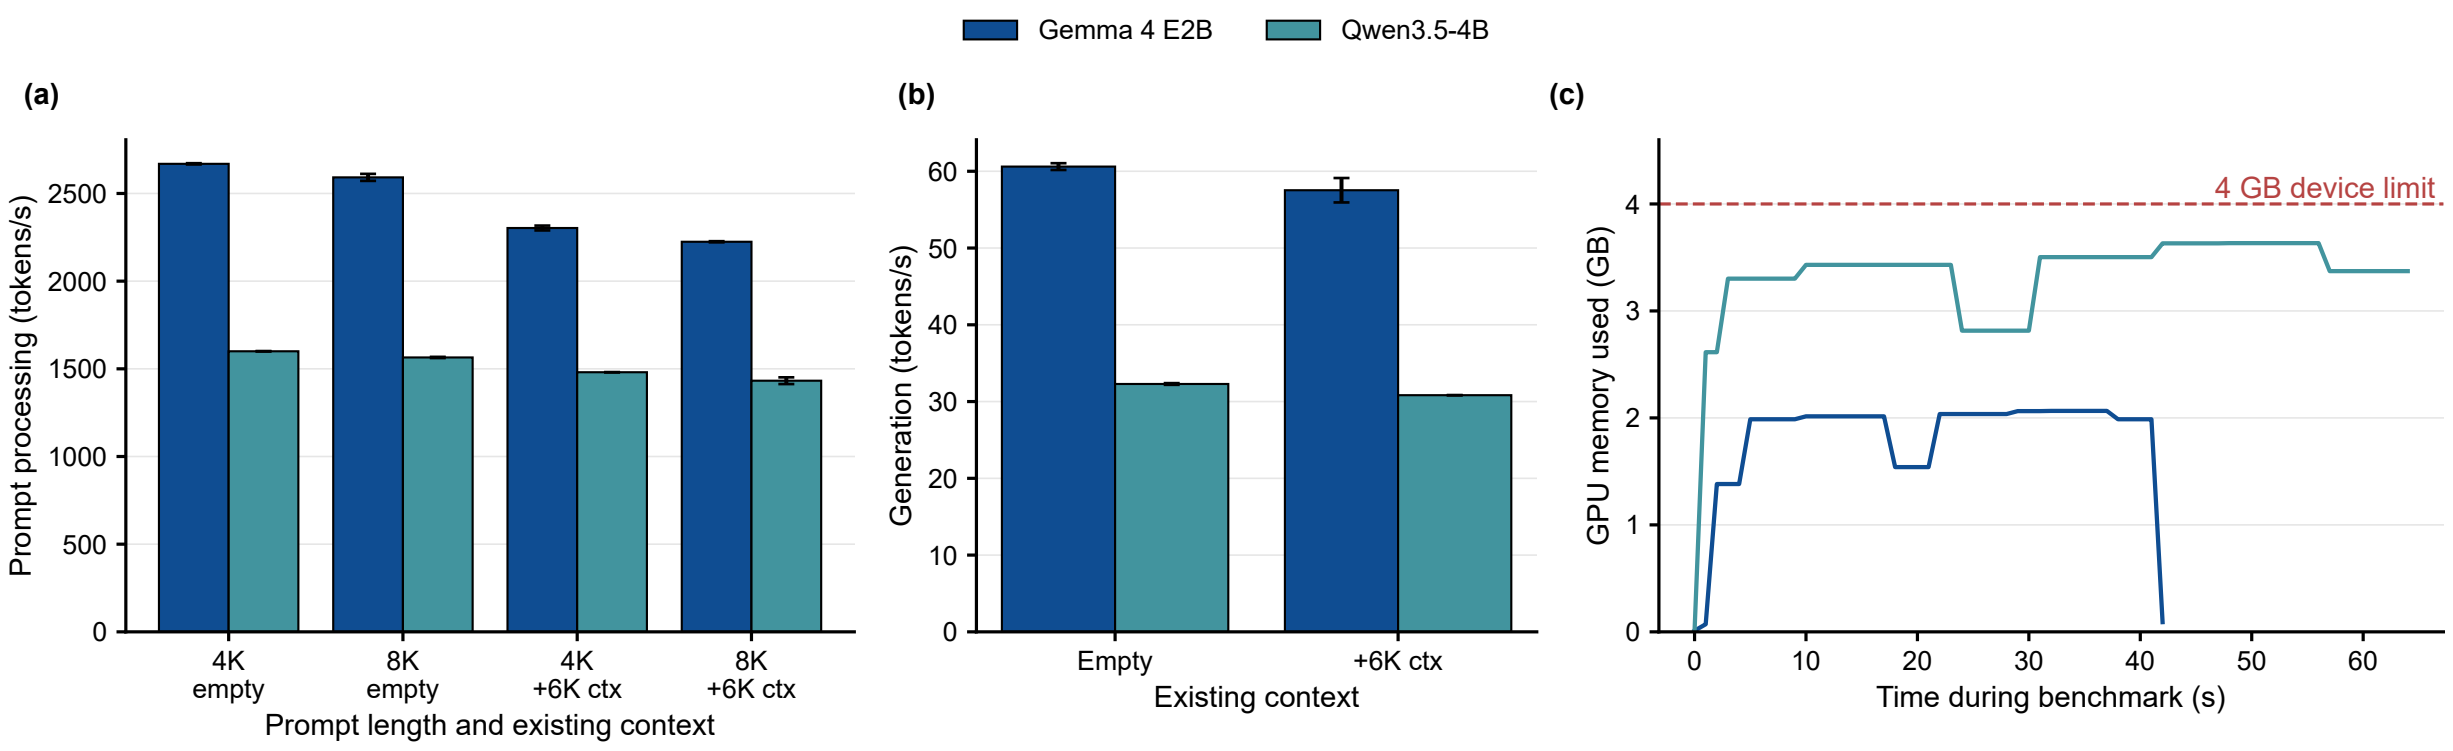
Gemma 4 E2B
Qwen3.5-4B
(a)
Prompt processing (tokens/s)
0
500
1000
1500
2000
2500
4K empty
8K empty
4K +6K ctx
8K +6K ctx
Prompt length and existing context
(b)
Generation (tokens/s)
0
10
20
30
40
50
60
Empty
+6K ctx
Existing context
(c)
GPU memory used (GB)
4 GB device limit
0
1
2
3
4
0
10
20
30
40
50
60
Time during benchmark (s)