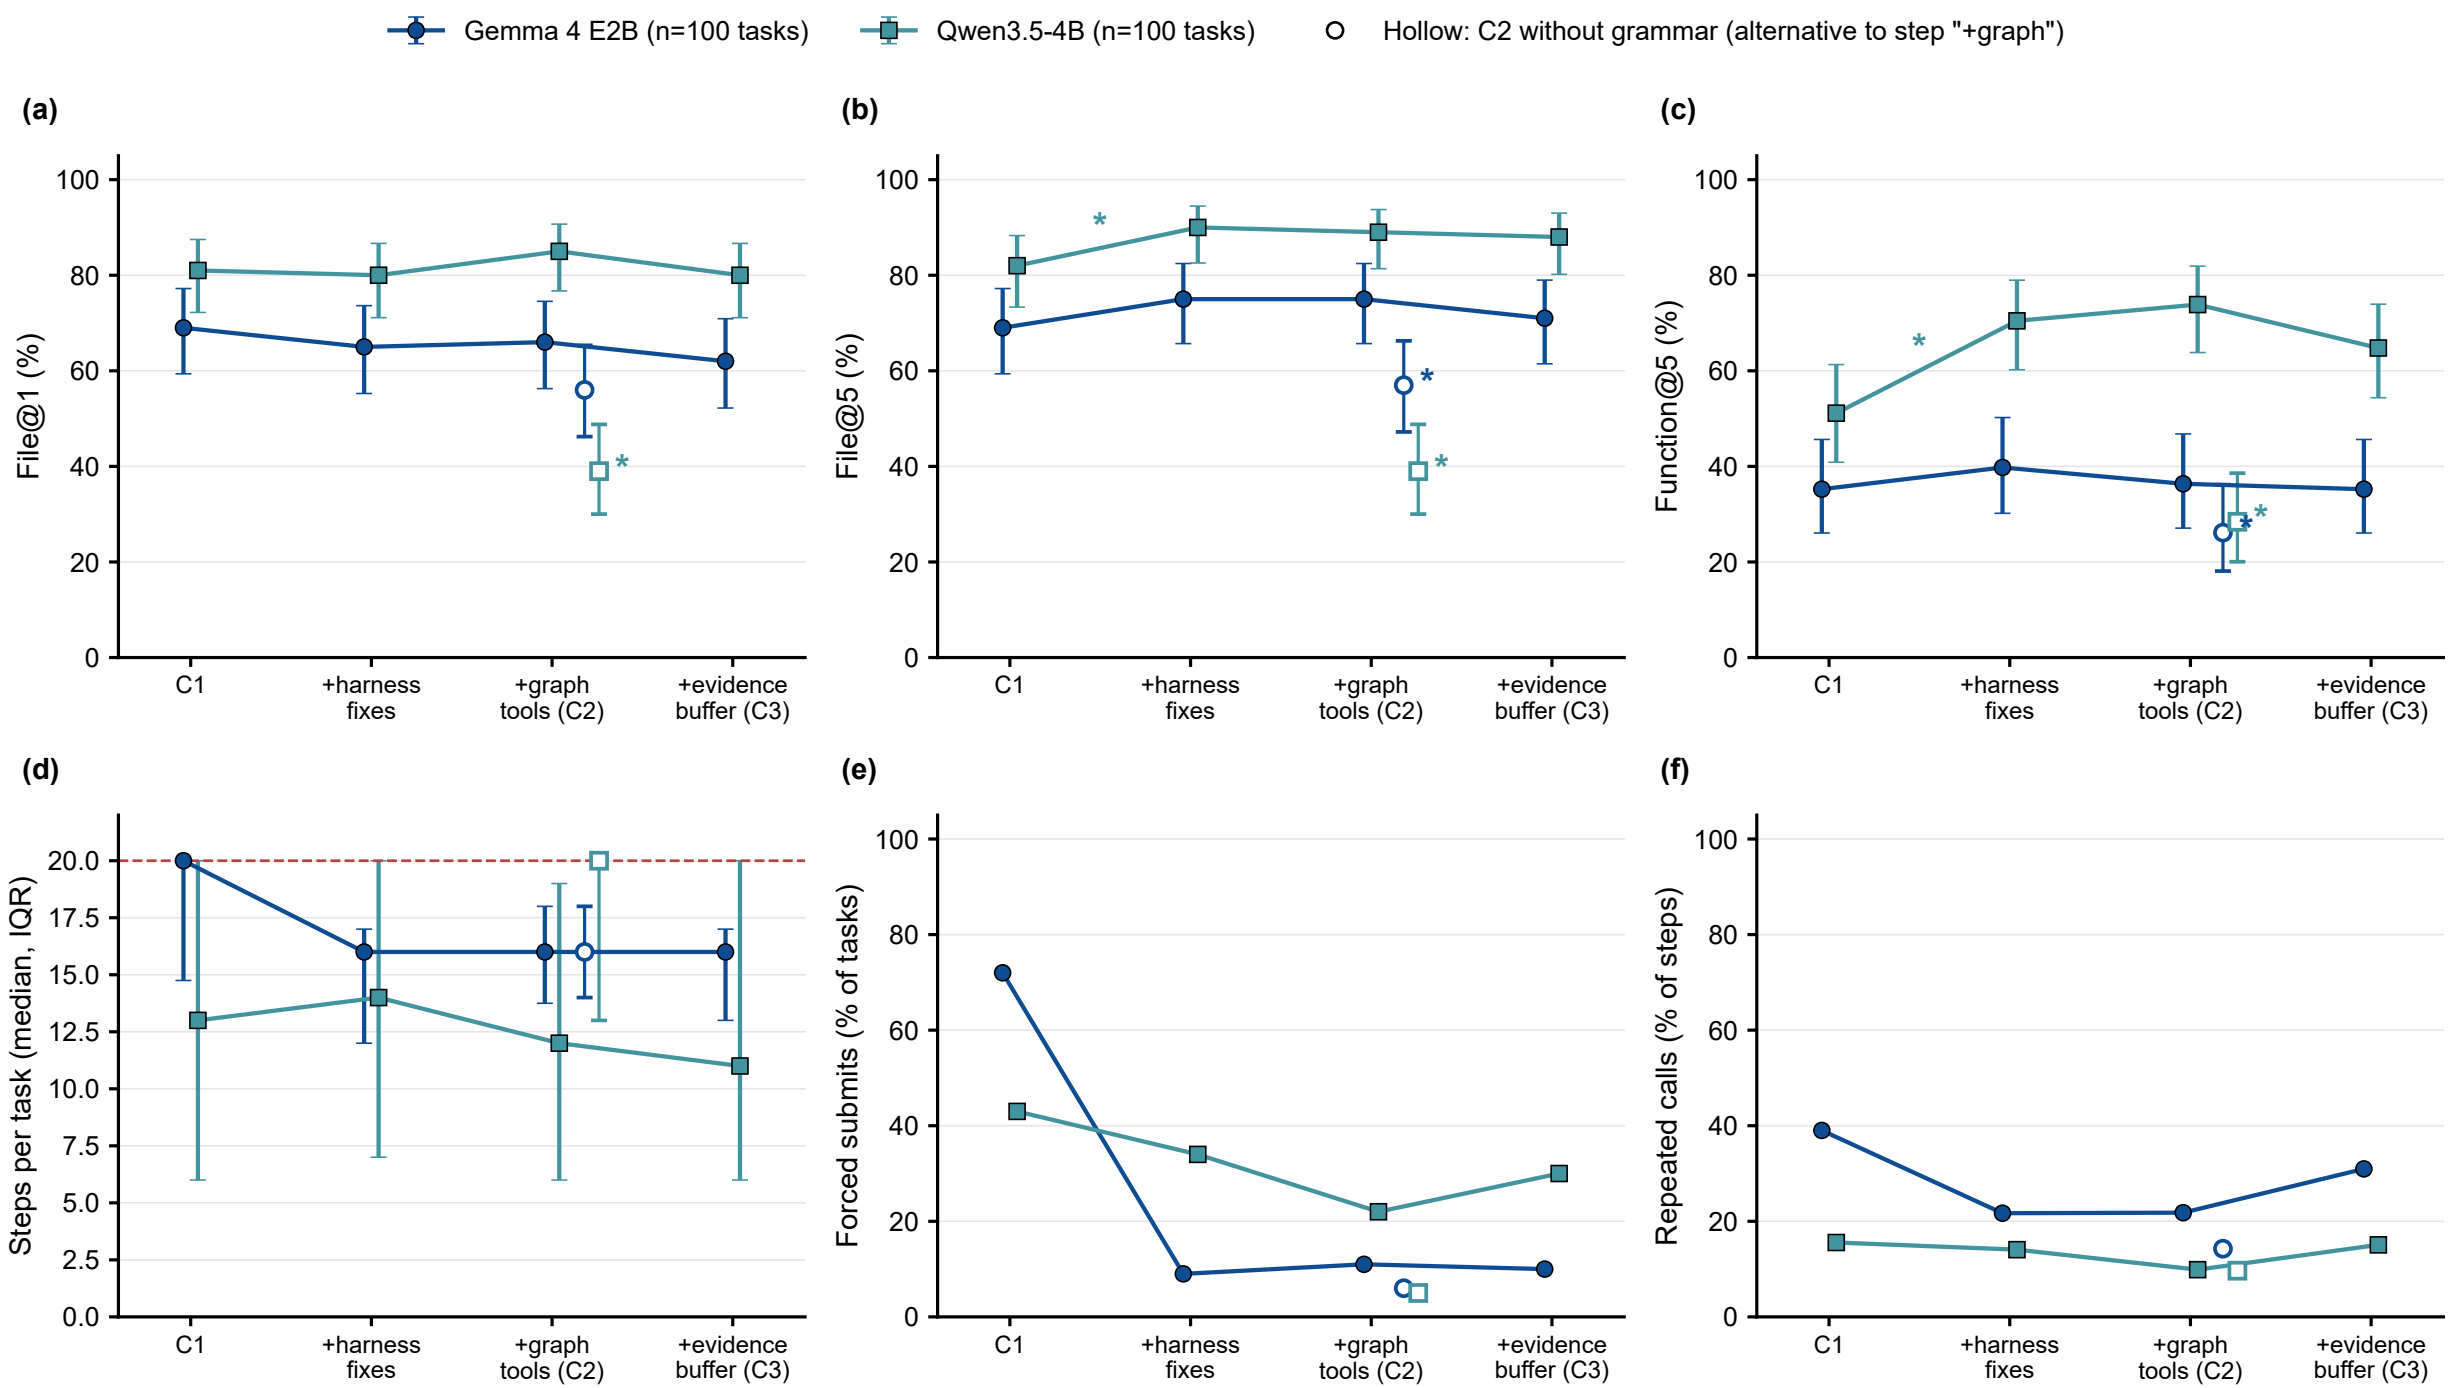

Each configuration adds one change to the one on its left; all run on the same tasks. Error bars: 95% Wilson intervals (a-c), interquartile range (d).
* exact McNemar test on paired per-task outcomes against the previous configuration, $p < 0.05$ (the hollow alternative is tested against the same previous configuration).